

she codes;

GLOBAL ACTIVISM

Eti Mayer
Final Project – NOV 2021



ETI MAYER

- Junior Data Scientist – Certified DS from Bar Ilan
- DS Skills: ML Models, R, Python (DS packages), SQL Server
- Financial back office in an activist hedge fund –



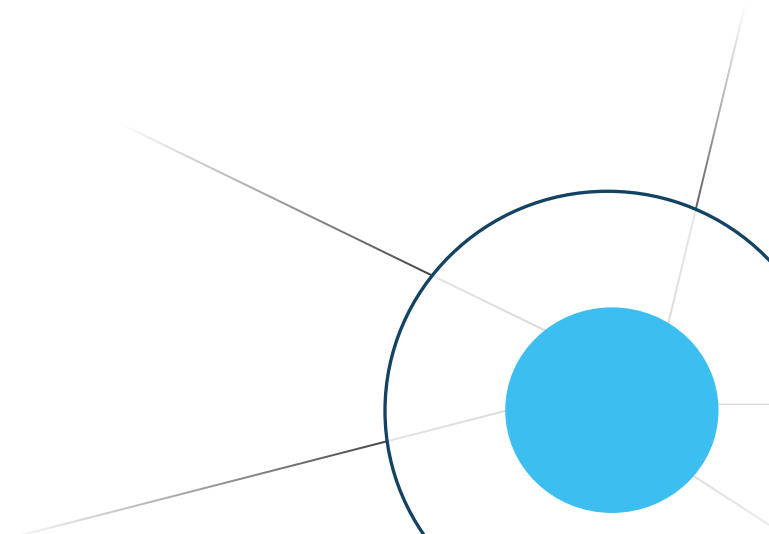
Experience:

- She Code Final Project – Global Activism
- Israeli Security Agency Hackathon 2021 – DS Team Lead
- Final project in the DS course – NoShow dataset



GLOBAL ACTIVISM

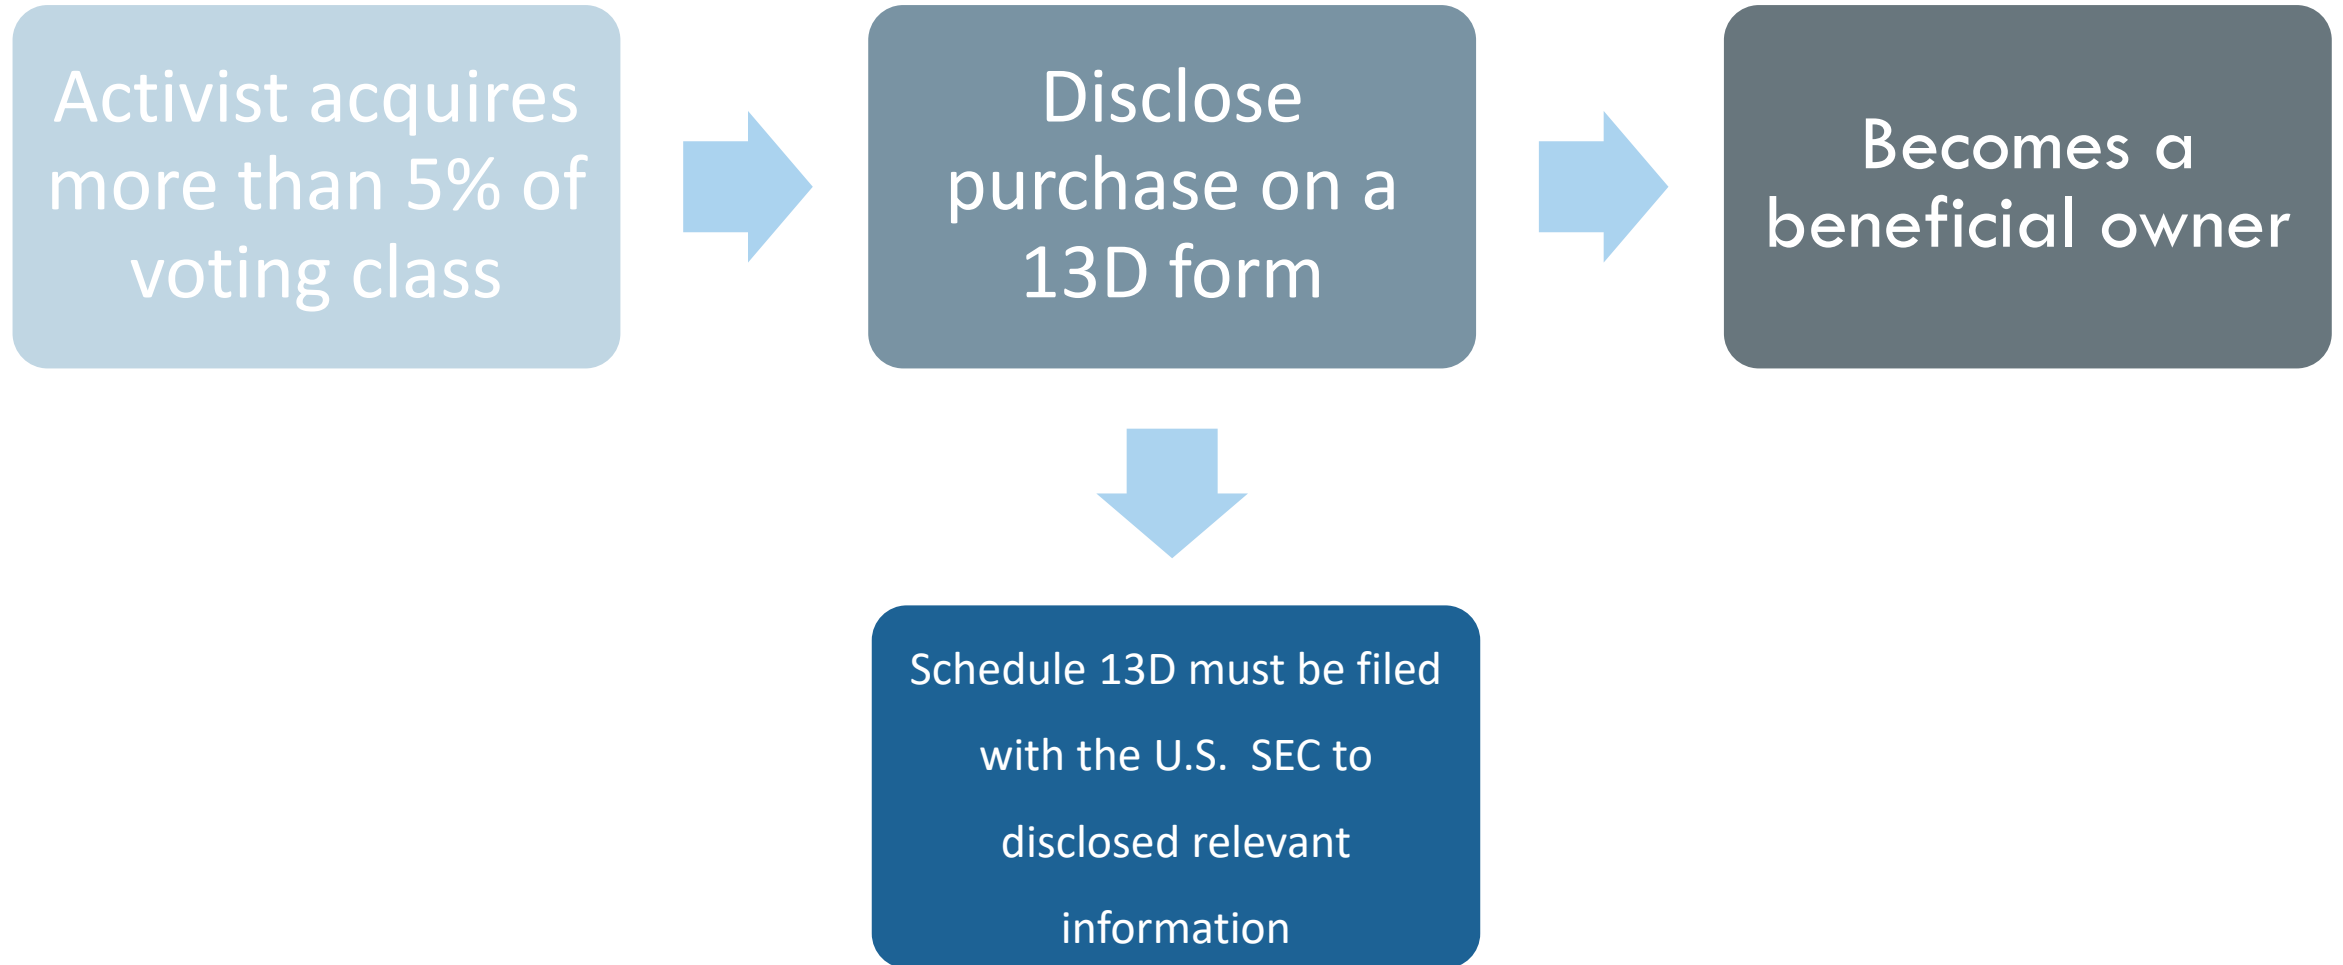
- Domain Field: Global Capital Market
- Main Project Purpose: Predict Activists' Campaign Succession for an Activism Follower Fund
- Data Source: Purchased Database.
+ Edgar SEC database



WHAT IS AN ACTIVIST CAMPAIGN?

- An activist investor buys a significant stake in a public company
- **Strategy:** Influence on how the company is run, e.g.:
 1. Obtaining seats on BOD.
 2. Taking private
 3. Etc.
- **Objective:** Creating value to share holders:
 - >50% net IRR over ref index, S&P500

● HOW TO FIND ACTIVISTS' CAMPAIGNS?



GLOBAL SUCCESS



Aims	Examples
Strategic Change / Forced Sale	       
ROI	     
Assets Sale / Reconciliation of Investments	     
Impact on Merging	      
Strategic Change	      
Replacing Management / Corporate Governance	      

THE DATA

```
In [3]: 1 con <- dbConnect(odbc(),
2           Driver = "SQL Server",
3           Server = "DESKTOP-AAGNMGA\\SQLEXPRESS",
4           Database = "Activism",
5           Trusted_Connection = "True")
```

```
In [4]: 1 df <- dbReadTable(con,"activist_holdings_v_SC")
2       head(df)
```

A data.frame: 6 × 189

	Investor.ID	Activist	ActivistHQ	ActivistRegion	Founded	FirstDateInvestedByActivisit	CurrentHolding	StatusCurrent	StatusExisted	DateExited	...	Seal
	<int>	<chr>	<chr>	<chr>	<int>	<chr>	<dbl>	<int>	<int>	<chr>	...	
1	2	Aberdeen Asset Management PLC	UK	WestEurope	1983	2017-01-03	NA	1	0	NA	...	
2	2	Aberdeen Asset Management PLC	UK	WestEurope	1983	2010-09-10	11	1	0	NA	...	
3	2	Aberdeen Asset Management PLC	UK	WestEurope	1983	2015-09-17	NA	0	1	2016-10-10	...	

Aberdeen

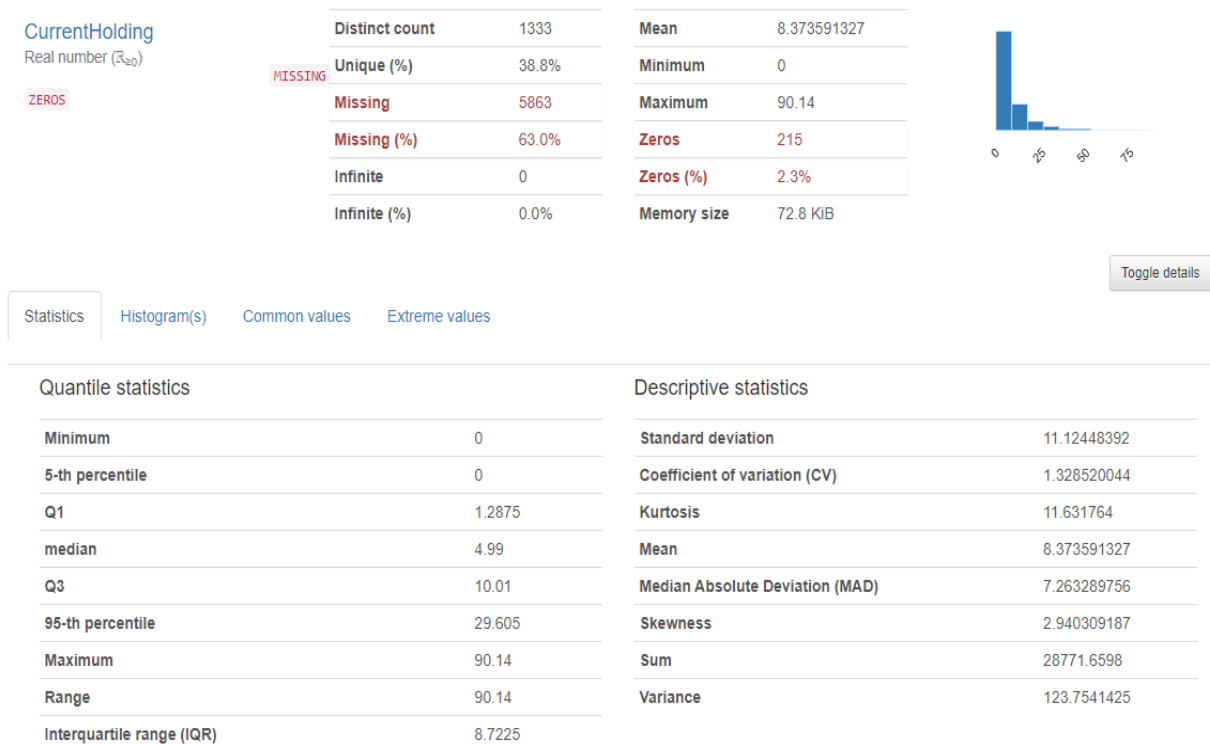
A Data.Frame: 9299x189 ; **Target=Succession (Categorical 1/0)**

Target: CASE WHEN (α.[Follower Return Annualised (%)]/ α.[S&P Change Annualised (%)]-1>0.5)
THEN (1) ELSE (0) END AS Succession



EDA — Data Visualization

- Using Pandas-Profiling package for statistical and visualized insights.
- ## Full Report



EDA - Correlation

- Correlation matrix between continues variables

```


1 res <- NULL
2
3 for(i in names(Activism_continuousV)) {
4   rw <- NULL
5   for(j in names(Activism_continuousV)) {
6     rw <- cbind(rw, cor.test(x=df[[i]], y=df[[j]], method="spearman")$estimate)
7   }
8   res <- rbind(res, rw)
9 }
10 row.names(res) <- names(Activism_continuousV)
11 colnames(res) <- names(Activism_continuousV)

```

```

1 corrplot(res, method = "shade", type = "upper", is.corr = TRUE, tl.cex=1.5, tl.col = "Blue", tl.srt = 90, na.rm = TRUE)
2 options(repr.plot.width = 10, repr.plot.height = 10)
3

```





EDA - Correlation

- Correlation between target 'Succession' to categorical variables

```


1 def conditional_entropy(x,y):
2     # entropy of x given y
3     y_counter = Counter(y)
4     xy_counter = Counter(list(zip(x,y)))
5     total_occurrences = sum(y_counter.values())
6     entropy = 0
7     for xy in xy_counter.keys():
8         p_xy = xy_counter[xy] / total_occurrences
9         p_y = y_counter[xy[1]] / total_occurrences
10        entropy += p_xy * math.log(p_xy/p_y)
11    return entropy
12
13 def theil_u(x,y):
14     s_xy = conditional_entropy(x,y)
15     x_counter = Counter(x)
16     total_occurrences = sum(x_counter.values())
17     p_x = list(map(lambda n: n/total_occurrences, x_counter.values()))
18     s_x = ss.entropy(p_x)
19     if s_x == 0:
20         return 1
21     else:
22         return (s_x - s_xy) / s_x

```

```

1 theilu = pd.DataFrame(index=['Succession'],columns=df.columns)
2 columns = df.columns
3 for j in range(0,len(columns)):
4     u = theil_u(df['Succession'].tolist(),df[columns[j]].tolist())
5     theilu.loc[:,columns[j]] = u
6 theilu.fillna(value=np.nan,inplace=True)
7 plt.figure(figsize=(20,5))
8 sns.heatmap(theilu,annot=True,fmt='.2f')
9 plt.show()

```



EDA — Data Cleaning - Outliers

Detecting Outliers
using outliers' matrix




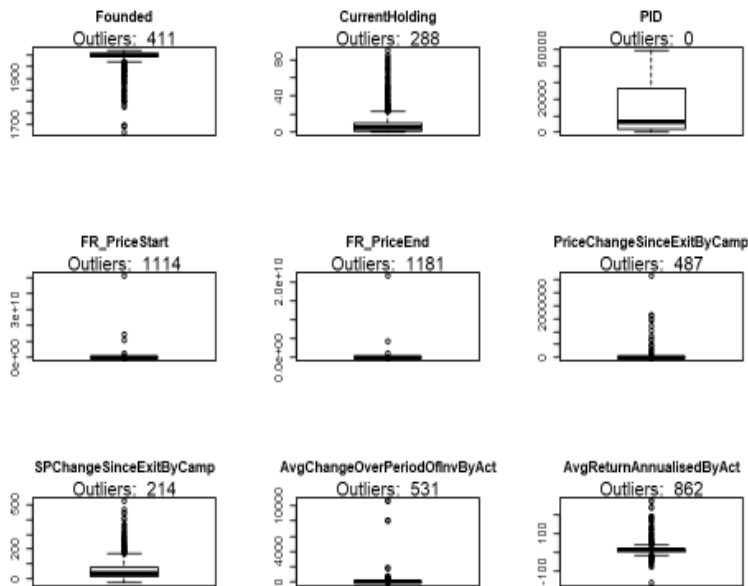
Checking the
distribution
with(blue) / without
outliers(red)



Correlation matrix -
Y + X(with outliers)
& X(without outliers)



Result DF



EDA — Data Cleaning - Outliers

Detecting Outliers
using outliers' matrix



Checking the
distribution
with(blue) / without
outliers(red)




Correlation matrix -
Y + X(with outliers)
& X(without outliers)



Result DF

	with <dbl>	pv_w <dbl>	without <dbl>	pv_wo <dbl>	diff <dbl>	cor.drop <chr>	dis.drop <chr>	drop <fct>
Founded	0.02447008	3.397179e-02	-0.002885733	8.078509e-01	0.88207095	No	Yes	Drop
CurrentHolding	-0.20498188	9.742106e-34	-0.183220686	4.694388e-25	0.10616154	Yes	Yes	Drop
FR_PriceStart	0.08316077	2.335499e-15	0.107911859	5.052674e-22	-0.29762941	Yes	Yes	Drop
FR_PriceEnd	0.38664775	2.146221e-320	0.414456884	0.000000e+00	-0.07192369	Yes	Yes	Drop
PriceChangeSinceExitByCamp	0.02817042	6.829174e-02	0.113201991	4.438121e-12	-3.01847070	No	Yes	Drop
SPChangeSinceExitByCamp	0.08762148	2.183701e-07	0.077937526	7.503105e-06	0.11052034	Yes	Yes	Drop
AvgChangeOverPeriodOfInvByAct	0.09537587	8.421008e-15	0.060024133	2.834695e-06	0.37065707	No	No	Leave
AvgReturnAnnualisedByAct	0.25382696	1.808149e-97	0.194388531	4.861968e-50	0.23416910	Yes	No	Drop
AvgSPReturnByAct	-0.20973284	1.891091e-66	-0.195151021	6.636532e-56	0.06952569	Yes	No	Drop



EDA — Data Cleaning — Missing Values

- Summarizing missing values after outliers' treatment

```
1 getMissingness(df_noout)
```

\$missingness
A data.frame: 116 × 3

	var	na_count	rate
	<fct>	<dbl>	<dbl>
	AH_PruchaseValueMnLocal	8385	90.2
	AvgAnnualisedReturnByMC	8329	89.6
	AH_PricePerShareLocal	8220	88.4
	Buyer	7866	84.6
	ThirteenD	6845	73.6
	SeatsProposed	6751	72.6
	SeatsGained	6725	72.3
	CurrentHolding	6151	66.1


DelayShareholderMeeting	7	0.1
CloseFund	7	0.1
ClosedAGM	7	0.1
ActivistIssuesPublicLetter	7	0.1
ActivistLetterToRegulatoryBodies	7	0.1
LitigationInitiated	7	0.1
ConsentSolicitationInitiated	7	0.1
SECFiling	7	0.1
AvgSPReturnByIND	2	0.0

\$message
'This dataset has 0 (0%) complete rows. Original data has 9299 rows.'

\$rows
NULL

EDA — Data Cleaning — Missing Values


- Heatmap for missing values



EDA — Data Cleaning — Missing Values

- Determine the Missingness Generation Mechanism

By the distribution and the t-test I assume that the missing mechanism is MNAR



```
[1] "SPChangeSinceExitByCamp by NA PositiveReturnByIND"
```

```
Welch Two Sample t-test
```

```
data: b1[[i]] and b2[[i]]  
t = -0.4457, df = 1153.3, p-value = 0.6559
```

EDA — Data Cleaning — Missing Values

- Treatment:
 1. Quantile numeric variable and adding 'miss' category

```
1 # quantile numeric variable
2 for (i in intVar){
3   df.na[[i]] <- as.factor(ifelse (is.na(df.na[[i]]), 'miss', quantile(df.na[[i]], probs = seq(0, 1, 0.25), na.rm = TRUE,
4 }
5 getMissingness(df.na)
```

2. For categorical vars with $>10\%$ missing – add 'miss'

```
var na_count rate
<fct> <dbl> <dbl>
```

```
$message
```

```
'This dataset has 9299 (100%) complete rows. Original data has 9299 rows.'
```

```
$rows
```

```
NULL
```

Re-Check for Outliers and missing values




EDA — Notes and challenges

- A wrong 'ifelse' transcript converted big part of the data to NA.
- Outliers' treatment

CLUSTER ANALYSIS

- Adding 2 type of clusters to the data, Threshold 3 and 4 :
 1. Hierarchical clustering – hclust.



2. Gaussian mixture models – mclust.

```
1 library(mclust)
2
3 ### Mclust: implementation of gaussian mixture model
4
5 mcl_model3 <- Mclust(df, 3)
6
7 summary(mcl_model3)
8
9 mcl_model4 <- Mclust(df, 4)
10
11 summary(mcl_model4)
```

FEATURE SELECTION STRATEGY

- 2 methods for feature selection:
 1. Univariable Analysis – Selecting features with $p_value \leq 0.05$
 2. Multivariable Analysis – Using Lasso, Ridge, Naïve Bayse, Random forest, etc.
- Summarization and Selection of Variables – Threshold=1

```
1 varSel['Sum'] = np.sum(varSel,axis=1)
2 varSel
```

	Variable	Univariable	Lasso	Ridge	ComplementNB	RandomForest	CART	GradientBoost	ADABOOST	SVM	Sum
0	ActivistHQ_1	0	0	0	0	1	1	0	0	0	2
1	ActivistHQ_2	1	0	0	0	1	1	1	0	0	4
2	ActivistHQ_3	1	0	0	0	1	0	0	0	0	2
3	ActivistHQ_4	0	0	1	0	0	0	0	0	0	1
4	ActivistHQ_5	1	0	0	0	1	0	0	0	0	2
...
3280	Unresolved_2	1	0	0	0	1	1	0	0	0	3
3281	hclust3	1	0	0	0	1	1	0	0	0	3
3282	hclust4	1	0	0	0	1	1	0	0	1	4
3283	Mclust3	1	0	0	0	1	1	1	0	0	4
3284	Mclust4	1	0	0	0	1	1	1	0	1	5

3285 rows x 11 columns



● FEATURE SELECTION — Notes and challenges

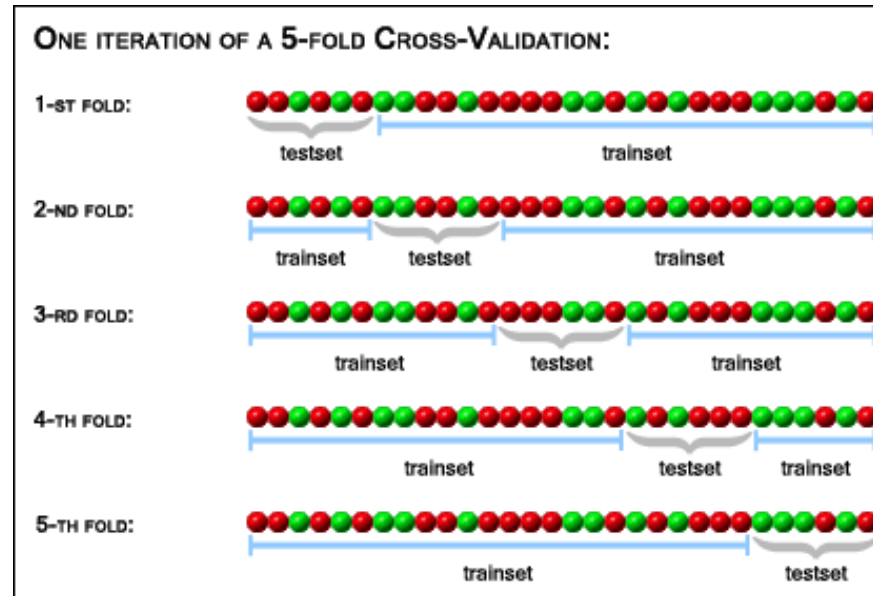
- Memory error due to high cardinality on 'one hot encoding'.
- Using data batches for the encoder gave me different features which I couldn't join back to DF

DATA PRE-PROCCING — Preparing the data

- Splitting data to train-dev-test

You got a perfectly balanced training and test datasets

- Using cross validation for Train-Dev




CLASSIFICATION MODELS — Metric Selection

- I chose 'Sensitivity' score with FPR – The Gain driven approach

The false positive rate is $FPR = \frac{FP}{FP+TN} = 1 - \text{Specificity}$

Because I wanted to make sure I won't miss any campaign that might be successful.

Sensitivity vs. Specificity





CLASSIFICATION MODELS - CV strategy Selection

- **StratifiedKFold** and **StratifiedShuffleSplit** for imbalance data

Category	#
0	6,979
1	2,320

Imbalance Ratio: 1 / 4



CLASSIFICATION MODELS - Results

- Model with the highest sensitivity score and the lowest FPR



CLASSIFICATION MODELS



	Name	Model	AUC-train	AUC-test	AUC_diff	F1-train	F1-test	F1_diff	Sensitivity-train	Sensitivity-test	sens_diff	fpr-train	fpr-test	fpr_diff
2	SVC	mod3	0.847208	0.818621	0.028588	0.857845	0.857757	0.000087	1.000000	1.000000	0.000000	1.000000	1.000000	0.000000
0	SVC	mod3	0.841789	0.812042	0.029747	0.857801	0.857845	0.000044	1.000000	1.000000	0.000000	1.000000	1.000000	0.000000
1	SVC	mod3	0.845187	0.824536	0.020651	0.857801	0.857845	0.000044	1.000000	1.000000	0.000000	1.000000	1.000000	0.000000
1	Gradient Boosting Classifier	mod8	0.897174	0.849202	0.047971	0.894374	0.879195	0.015180	0.953455	0.937724	0.015731	0.538877	0.589633	0.050756
2	Gradient Boosting Classifier	mod8	0.895010	0.851578	0.043432	0.891282	0.876722	0.014561	0.947745	0.934814	0.012931	0.539957	0.596112	0.056156
2	Logistic Regression	mod1	0.897094	0.848336	0.048758	0.896873	0.874362	0.024511	0.941661	0.919771	0.021890	0.463283	0.555076	0.091793
1	Logistic Regression	mod1	0.897233	0.844349	0.052885	0.896694	0.877349	0.019345	0.936985	0.919112	0.017873	0.461123	0.531317	0.070194
0	Gradient Boosting Classifier	mod8	0.898423	0.846752	0.051671	0.891411	0.869654	0.021757	0.941998	0.916965	0.025033	0.517279	0.578834	0.061555
0	Logistic Regression	mod1	0.898274	0.841868	0.056406	0.899092	0.873326	0.025766	0.939492	0.910523	0.028969	0.453564	0.526996	0.073434
1	ADABOOST	mod6	0.871375	0.841471	0.029905	0.879447	0.867238	0.012209	0.911565	0.904796	0.006769	0.487041	0.548596	0.061555
2	Random Forest Classifier	mod5	0.999095	0.760225	0.238869	0.992342	0.852991	0.139351	0.997137	0.904011	0.093125	0.037797	0.650108	0.612311
1	XGboost	mod7	0.981824	0.833186	0.148637	0.956293	0.864902	0.091391	0.975295	0.900501	0.074794	0.194384	0.548596	0.354212
0	ADABOOST	mod6	0.870145	0.826714	0.043431	0.880318	0.864436	0.015882	0.911207	0.896922	0.014285	0.479482	0.537797	0.058315
0	XGboost	mod7	0.985828	0.833010	0.152818	0.963237	0.868330	0.094907	0.980308	0.896922	0.083386	0.166307	0.509719	0.343413
2	ADABOOST	mod6	0.871031	0.836379	0.034652	0.879227	0.866136	0.013091	0.911954	0.896848	0.015106	0.490281	0.524838	0.034557
1	Random Forest Classifier	mod5	0.999260	0.752747	0.246513	0.993022	0.849437	0.143585	0.993555	0.890480	0.103076	0.022678	0.622030	0.599352
0	Random Forest Classifier	mod5	0.999175	0.773250	0.225925	0.992123	0.849551	0.142572	0.992123	0.881174	0.110949	0.023758	0.583153	0.559395
2	XGboost	mod7	0.982256	0.834246	0.148009	0.960434	0.856745	0.103690	0.981747	0.880372	0.101374	0.168985	0.526996	0.338013
2	Decision Tree Classifier	mod4	0.999988	0.693839	0.306149	0.998586	0.853901	0.144686	0.997137	0.862464	0.134673	0.000000	0.475162	0.475162
0	Decision Tree Classifier	mod4	0.999984	0.678691	0.321294	0.998386	0.842781	0.155605	0.996778	0.846099	0.150679	0.000000	0.488121	0.488121
1	Decision Tree Classifier	mod4	0.999993	0.662455	0.337537	0.998925	0.832200	0.166725	0.998210	0.832498	0.165712	0.001080	0.507559	0.506479
1	ComplementNB	mod2	0.757936	0.681560	0.076376	0.789504	0.760886	0.028618	0.732546	0.712956	0.019589	0.371490	0.485961	0.114471
2	ComplementNB	mod2	0.759515	0.717470	0.042045	0.794563	0.766847	0.027715	0.742663	0.709169	0.033494	0.382289	0.423326	0.041037
0	ComplementNB	mod2	0.761277	0.695748	0.065529	0.796791	0.765163	0.031628	0.746867	0.704366	0.042501	0.385529	0.412527	0.026998



CLASSIFICATION MODELS — Model selected

Decision Tree Classifier is the best considering both parameters

	Name	Model	AUC-train	AUC-test	AUC_diff	F1-train	F1-test	F1_diff	Sensitivity-train	Sensitivity-test	sens_diff	fpr-train	fpr-test	fpr_diff
2	Decision Tree Classifier	mod4	0.999988	0.693839	0.306149	0.998566	0.853901	0.144666	0.997137	0.862464	0.134673	0.000000	0.475162	0.475162
0	Decision Tree Classifier	mod4	0.999984	0.678691	0.321294	0.998386	0.842781	0.155605	0.996778	0.846099	0.150679	0.000000	0.488121	0.488121
1	Decision Tree Classifier	mod4	0.999993	0.662455	0.337537	0.998925	0.832200	0.166725	0.998210	0.832498	0.165712	0.001080	0.507559	0.506479



CLASSIFICATION MODELS — Fine Tuning

- Random search and Grid Search for hyperparameter optimization

Random search

```
1 # The function to measure the quality of a split
2 criterion = ['gini', 'entropy']
3 # The strategy used to choose the split at each node.
4 # Supported strategies are "best" to choose the best split and "random" to choose the best random split
5 splitter = ['best', 'random']
6 # The minimum number of samples required to split an internal node
7 min_samples_split = [2,4,8,16,32,40]
8 # The minimum number of samples required to be at a leaf node
9 min_samples_leaf = [2,4,6,8,10,12,14,16,18,20]
10
11 random_grid = {'criterion': criterion,
12                'splitter': splitter,
13                'min_samples_split': min_samples_split,
14                'min_samples_leaf': min_samples_leaf}
15
16 print(random_grid)
```

```
{'criterion': ['gini', 'entropy'], 'splitter': ['best', 'random'], 'min_samples_split': [2, 4, 8, 16, 32, 40], 'min_samples_leaf': [2, 4, 6, 8, 10, 12, 14, 16, 18, 20]}
```

```
1 DTC_random.best_params_
{'splitter': 'best',
 'min_samples_split': 40,
 'min_samples_leaf': 18,
 'criterion': 'gini'}
```



CLASSIFICATION MODELS — Fine Tuning

- Grid search - We decide which parameters and how (not randomly)

```
: 1 # Create the parameter grid
2 param_grid = {'splitter': ['best'],
3               'min_samples_split': [35,40,45],
4               'min_samples_leaf': [17,18,19],
5               'criterion': ['gini']}
6 }
7
8 grid_search = GridSearchCV(estimator = DTC, param_grid = param_grid, cv = sss,
9                             verbose=2, n_jobs=-1)
```

```
1 grid_search.best_params_
```

```
{'criterion': 'gini',
 'min_samples_leaf': 18,
 'min_samples_split': 40,
 'splitter': 'best'}
```

	Base Score	Random Score	Grid Score
0	0.847222	0.891762	0.966696
1	0.875000	0.902778	0.961368
2	0.863027	0.867816	0.962256

We can see the improvement between base, random and grid fine tuning

PREDICTIONS

```
1 # make a prediction
2 y_pred = best_grid.predict(X_test)
3
```

```
1 pd.crosstab(y_test,y_pred)
```

```
C:\Users\Mayer\Anaconda3\lib\site-packages\pandas\core\series.py:100:
Warning: The following object is a categorical object instead of an ndarray :
vec = libmissing.isnaobj_old(values.ravel())
```

	col_0	0	1
Succession			
0	754	635	
1	322	3868	

```
1 yprob = best_grid.predict_proba(X_test)
2 yprob = pd.DataFrame(yprob)
3 yprob
```

	0	1
0	0.289474	0.710526
1	0.090909	0.909091
2	0.000000	1.000000
3	0.000000	1.000000
4	0.000000	1.000000
...
5574	0.000000	1.000000
5575	0.217391	0.782609
5576	0.000000	1.000000
5577	0.916667	0.083333

FUTURE DEVELOPMENT

- Training day0 set (features from 1 3D)
- Selecting model by a conservative approach (Specificity)
- Production

The background is a complex network of thin grey lines connecting various circular nodes. The nodes are in different colors: dark blue, light blue, and grey. Some nodes are larger than others, and some are surrounded by concentric circles. The overall effect is a sense of interconnectedness and digital communication.

THANK YOU